

UNIT 3 STRUCTURED QUERY LANGUAGE

Structure

- 3.0 Introduction
- 3.1 Objectives
- 3.2 Categories of SQL Commands
- 3.3 Data Definition
- 3.4 Data Manipulation Statements
 - 3.4.1 SELECT - The Basic Form
 - 3.4.2 Subqueries
 - 3.4.3 Functions
 - 3.4.4 GROUP BY Feature
 - 3.4.5 Updating the Database
 - 3.4.6 Data Definition Facilities
- 3.5 Views
- 3.6 Summary
- 3.7 Further Reading

3.0 INTRODUCTION

SQL is an acronym for Structured Query Language. It is available in a number of data base management packages based on the relational model of data, for example, in DB2 of the IBM and UNIFY of the UNIFY corporation.

Originally defined by D.D. Chamberlain in 1974, SQL underwent a number of modifications over the years. Today, SQL has become an official ANSI standard.

It allows for data definition, manipulation and data control for a relational database. The data definition facilities of SQL permit the definition of relations and of various alternative views of relations. Further, the data control facility gives features for one user to authorize other users to access his data. This facility also permits assertions to be made about data integrity. All the three major facilities of SQL, namely, data manipulation, data definition and data control are bound together in one integrated language framework.

3.1 OBJECTIVES

After going through this unit you will be able to:

- Differentiate SQL commands
- List data manipulation commands
- List data definition commands
- Make queries using data manipulation commands.

3.2 CATEGORIES OF SQL COMMANDS

SQL commands can be roughly divided into three major categories with regard to their functionality. Firstly, there are those used to create and maintain the database structure. The second category includes those commands that manipulate the data in such structures, and thirdly there are those that control the use of the database. To have all this functionality in a single language is a clear advantage over many other systems straightway, and must certainly contribute largely to the rumour of it being easy to use.

It's worth naming these three fundamental types of commands for future reference. Those that create and maintain the database are grouped into the class called DDL or Data Definition Language statements and those used to manipulate data in the tables, of which there are four, are the DML or Data Manipulation Language commands. To control usage of the data the DCL commands (Data Control Language) are used, and it is these three in conjunction plus one or two additions that define SQL. There are therefore no environmental

statements, as one finds so irritating in COBOL: for example, no statements to control program flow (if/then/else, perform, go to) and of course, no equivalent commands to open and close files, and read individual records. At this level then, it is easy to see where SQL gets its **end-user-tool** and **easy-to-use** tags.

The Data Definition Statements

To construct and administer the database there are two major DDL statements - **CREATE** and **DROP**, which form the backbone of many commands:

CREATE DATABASE to create a database **DROP DATABASE** to remove a database
CREATE TABLE to create a table **DROP TABLE** to drop a table **CREATE INDEX** to create an index on a column **DROP INDEX** to drop an index **CREATE VIEW** to create a view **DROP VIEW** to drop a view.

There may be some additional ones, such as **ALTER TABLE** or **MODIFY DATABASE**, which are vendor specific.

The Data Manipulation Statements

To manipulate data in tables directly or through views we use the four standard DML statements:

SELECT DELETE INSERT UPDATE

These statements are now universally accepted, as is their functionality, although the degree to which these commands support this functionality varies somewhat between products compare the functionality of different implementation of **UPDATE** for example.

Data Control

This deals with three issues

(a) Recovery and Concurrency

Concurrency is concerned with the manner in which multiple users operate upon the data base.

Each individual user can either reflect the updates of a transaction by using the **COMMIT** or can cancel all the updates of a transaction by using **ROLLBACK**.

(b) Security

Security has two aspects to it.

The first is the **VIEW** mechanism. A view of a relation can be created which hides the sensitive information and defines only that part of a relation which should be visible. A user can then be allowed to access this view.

CREATE VIEW LOCAL AS

SELECT * FROM SUPPLIER

WHERE SUPPLIER.CITY = 'Delhi'

The above view reveals only the suppliers of Delhi.

The second is by using **GRANT** operation. This shall grant one or more access rights to perform the data manipulative operations on the relations.

(c) Integrity Constraints

Integrity constraints are enforced by the system. For example, one can specify that an attribute of a relation will not take on null values.

3.3 DATA DEFINITION

Data definition in SQL is via the **create** statement. The statement can be used to create a table, index, or view (i.e., a virtual table based on existing tables). To create a table, the create statement specifies the name of the table and the names and data types of each column

of the table. Its format is :

```
create table <relation> (<attribute list>)
```

where the attribute list is specified as:

```
<attribute list> ::= <attribute name> (<data type>)[not null] <attribute list>
```

The data types supported by SQL depend on the particular implementation. However, the following data types are generally included: integer, decimal, real (i.e., floating point values), and character strings, both of fixed size and varying length. A number of ranges of values for the integer data type are generally supported, for example, integer and smallint. The decimal value declaration requires the specification of the total number of decimal digits for the value and (optionally), the number of digits to the right of the decimal point. The number of fractional decimal digits is assumed to be zero if only the total number of digits is specified.

```
<data type> ::= <integer> | <smallint> | <char(n)> | <float> | <decimal (p,q)>
```

In addition, some implementations can support additional data types such as bit strings, graphical strings, logical, data, and time. Some DBMSs support the concept of date. One possible implementation of date could be as eight unsigned decimal digits representing the data in the yyyyymmdd format. Here yyyy represents the year, mm represents the month and dd represents the day. Two dates can be compared to find the one that is larger and hence occurring later. The system ensures that only legal date values are inserted (19860536 for the date would be illegal) and functions are provided to perform operations such as adding a number of days to a date to come up with another date or subtracting a date from the current date to find the number of days, months, or years. Date constants are provided in either the format given above or as a character string in one of the following formats: mm/dd/yy; mm/dd/yyyy; dd-mmm-yy; dd-mmm-yyyy. In this text we represent a date constant as eight unsigned decimal digits in the format yyyyymmdd.

The employee relation for the hotel database can be defined using the create table statement given below. Here, the Empl_No is specified to be not null to disallow this unique identifier from having a null value. SQL supports the concept of null values and, unless a column is declared with the not null option, it could be assigned a null value.

create table	EMPLOYEE
Empl_No	integer not null,
Name	char (25),
Skill	char (20)
Pay-Rate	decimal (10,2))

The definition of an existing relation can be altered by using the alter statement. This statement allows a row column to be added to an existing relation. The existing tuples of the altered relation are logically considered to be assigned the null value for the added column. The physical alteration occurs to a tuple only during an update of the record.

```
alter table existing-table-name  

add column-name data-type [....]  

alter table EMPLOYEE  

add Phone_Number decimal (10)
```

The create index statement allows the creation of an index for an already existing relation. The columns to be used in the generation of the index are also specified. The index is named and the ordering for each column used in the index can be specified as either ascending or descending. The cluster option could be specified to indicate that the records are to be placed in physical proximity to each other. The unique option specifies that only one record could exist at any time with a given value for the column(s) specified in the statement to create the index. (Even though this is just an access aid and a wrong place to declare the primary key). Such columns, for instance, could form the primary key of the relation and hence duplicate tuples are not allowed. One case is the ORDER relation where the key is the combination of the attribute Bill#, Dish#. In the case of an existing relation, an attempt to create an index with the unique option will not succeed if the relation does not satisfy this uniqueness criterion. The syntax of the create index statement is shown below:

```
create [unique] index name-of-index  

on existing-table-name
```

(column-name[ascending or descending]

[,column-name[order]....])

[cluster]

The following statement causes an index called empindex to be built on the columns Name and Pay_Rate. The entries in the index are ascending by Name value and descending by Pay_Rate. In this example there are no restrictions on the number of records with the same Name and Pay_Rate.

Create index empindex

on EMPLOYEE (Name asc, Pay_Rate desc);

An existing relation or index could be deleted from the database by the **drop** SQL statement.

The syntax of the drop statement is as follows:

drop table existing-table-name;

drop index existing-index-name;

3.4 DATA MANIPULATION

Data manipulation capabilities allows one to retrieve and modify contents of the data base. The most important of these is the **SELECT** operation which allows data to be retrieved from the data base.

The relation definitions that shall be used in the rest of the module are given below.

There are parts which are supplied by suppliers. S contains the details about each supplier. Turnover for a supplier is in terms of lakhs of rupees. Information regarding suppliers of specific parts is contained in SP whereas information about the parts themselves is contained in P.

S

S#	SNAME	SCITY	TURNOVER
10	CAUVERY	BANGALORE	50
11	NARMADA	BOMBAY	100
12	YAMUNA	DELHI	70
13	TAPI	BOMBAY	20

P

P#	WEIGHT	COLOUR	COST	SELLING PRICE
1	25	RED	10	30
2	30	BLUE	15	45
3	45	RED	20	45

SP

S#	P#	QTY
10	1	100
11	1	5
10	2	50
11	2	30
10	3	10
12	3	100
13	1	20

3.4.1 SELECT – The Basic Form

The select statement specifies the method of selecting the tuples of the relations(s). The tuples processed are from one or more relations specified by the form clause of the select statement.

The basic form of SELECT is

```
Select <target list>
from <relation list>
[where <predicate> ]
```

SELECT lists the attributes to be selected

FROM relations from which information is to be used

WHERE condition. The rows that qualify are those for which the condition evaluates to true.

Condition is a single predicate or a collection of predicates combined using the Boolean operators AND, OR and NOT.

The column names following SELECT are to be retrieved from the relations specified in the FROM part. WHERE specifies the condition that the tuples must satisfy in order to be part of the result.

Below we shall state first the retrieval query in English and they specify its SQL equivalent.

Unqualified Retrieval

- Get the part numbers of all the parts being supplied.

```
SELECT P#
FROM SP
```

P#
1
1
2
2
3
3
1

Part numbers getting repeated? That's right. SELECT does not eliminate duplicate rows (unlike the project operation of the relational algebra). In order to do that

- Get the part numbers of all the parts being supplied with no duplicates.

```
SELECT DISTINCT P#
FROM SP
```

P#
1
2
3

If all the columns of the relation are to be retrieved then one needn't list all of them. A * can be specified after SELECT to indicate retrieval of the entire relation.

- Get full details of all suppliers.

```
SELECT *
FROM S
```

S	S#	SNAME	SCITY	TURNOVER
	10	CAUVERY	BANGALORE	50
	11	NARMADA	BOMBAY	100
	12	YAMUNA	DELHI	70
	13	TAPI	BOMBAY	20

The ORDER BY clause

The result of a query can be ordered either in ascending (ASC) order or in descending (DESC) order.

- Get the supplier numbers and turnover in descending order of turnover.

```
SELECT S#, TURNOVER
FROM S
ORDER BY TURNOVER DESC
S#      TURNOVER
11      100
12      70
10      50
13      20
```

Instead of a column name, the ordinal position of the column in the result can be used. That is, the above query can be rewritten as

- SELECT S#, TURNOVER

```
FROM S
ORDER BY 2 DESC
```

The format of the order clause is

ORDER BY {int/col [ASC/DESC], }

col – column name

int – ordinal position of the column in the result table

ASC/DESC – Ascending or descending

If there is more than one specification, then the left-to-right specification corresponds to major-to-minor ordering. This is shown below.

- Get the supplier number and part number in ascending order of supplier number and descending order for the part supplied for each supplier.

```
SELECT S#, P#, QTY
FROM SP
ORDER BY S#, P# DESC
S#      P#      QTY
10      3       10
10      2       50
10      1       100
11      2       30
11      1       5
12      3       100
13      1       20
```

Qualified retrieval

The expression following WHERE specifies the condition that must be satisfied. Below we consider a few examples.

- Get the details of suppliers who operate from Bombay with turnover 50.

```
SELECT S.*
FROM S
WHERE CITY = 'BOMBAY' AND TURNOVER > 50
S#      SNAME      SCITY      TURNOVER
11      NARMADA    BOMBAY    100
```

The above form is a conjunction of comparison predicates. A comparison predicate is of the form

scalar-expr O scalar-expr

where O is any of the six relational operators

$=, < >, <, >, \leq, \geq$

and a scalar expression is an arithmetic expression with

operators as $+, -, *, /$

operands as col., function, constant

BETWEEN Predicate

- Get the part numbers weighing between 25 and 35

```
SELECT P#, WEIGHT
```

```
FROM P
```

```
WHERE WEIGHT BETWEEN 25 AND 35
```

P#	WEIGHT
1	25
2	30

The use of BETWEEN gives the range within which the values must lie. If the value should lie outside a range then BETWEEN is to be preceded by NOT. For example,

```
SELECT P#
```

```
FROM P
```

```
WHERE WEIGHT NOT BETWEEN 25 AND 35
```

P#	WEIGHT
3	45

would retrieve all part numbers whose weight is less than 25 or greater than 35 as shown above.

LIKE Predicate

This predicate is used for pattern matching. A column of type char can be compared with a string constant. The use of the word LIKE doesn't look for exact match but a form of wild string match. A % or - can appear in the string constant where

% stands for a sequence of n (≥ 0) characters

- stands for a single character

Examples

ADDRESS LIKE '%Bangalore%' - ADDRESS should have Bangalore somewhere as a part of it if the match is to succeed.

STRANGE STRING LIKE '%-%' ESCAPE'

Here, the normal meaning of - is overridden with the use of the escape character. STRANGE above will match with any string beginning with -

- Get the names and cities of suppliers whose name begin with C

```
SELECT SNAME, SCITY
```

```
FROM S
```

```
WHERE SNAME LIKE 'C%
```

```
SNAME SCITY
```

```
CAUVERY BANGALORE
```

When the data is to be retrieved from more than one relation then both the relation names is specified in the FROM clause and the join condition in the WHERE part.

10. For each part supplied, get part number and names of all cities supplying the part.

```
SELECT      P#, CITY
```

```
FROM        SP, S
```

```
WHERE       SP.S# = S.S#
```

```
P#      SCITY
```

```
1      BANGALORE
```

```
1      BOMBAY
```

```
2      BANGALORE
```

```
2      BOMBAY
```

```
3      BANGALORE
```

```
3      DELHI
```

```
1      BOMBAY
```

How does, then, one specify a join on the same relation?

11. Get pairs of supplier numbers such that both operate from the same city.

```
SELECT      FIRST.S#, SECOND.S#
```

```
FROM        S FIRST, S SECOND
```

```
WHERE       FIRST.CITY = SECOND.CITY
```

```
AND        FIRST.S# <>SECOND.S#
```

FIRST and SECOND are tuple variables, both ranging over S. The last line eliminates a supplier getting compared with himself.

```
S#      S#
```

```
11     13
```

```
13     11
```

But, we see that suppliers with numbers 11 and 13 are getting compared twice. Can that be avoided? How about < instead of <> ?

```
SELECT      FIRST.S#, SECOND.S#
```

```
FROM        S FIRST, S SECOND
```

```
WHERE       FIRST.CITY = SECOND.CITY
```

```
AND        FIRST.S# < SECOND.S#
```

Tests for NULL

An attribute can be tested for the presence or absence of null.

12. Get the supplier numbers whose turnover is null

```
SELECT      S#
```

```
FROM        S
```

```
WHERE       TURNOVER IS NULL
```

There is no tuple in the result of this query as in the sample

Can the last line in the above query be replaced by

```
WHERE      TURNOVER = NULL
```

Not really! It is incorrect as nothing (even NULL) is equal to NULL.

The format for specifying NULL is

col. ref IS [NOT] NULL

IN Predicate

This is to be used whenever you want to test whether an attribute value is one of a set of values. For example,

13. Get the part numbers that cost 20, 30 or 40 rupees.

```

SELECT      P.P#, SELLING PRICE
FROM        P
WHERE       SELLING PRICE IN (20, 40, 45)
P#          SELLING PRICE
2           45
3           45

```

It's a quicker way of specifying comparison.

The format of the predicate

scalar-expr [NOT] IN (atom list)

3.4.2 Subqueries

The expression following WHERE can be either a simple predicate as explained above or it can be a query itself! This part of the query following WHERE is called a Subquery.

A subquery, which in turn is a query, can have its own subquery and the process of specifying subqueries can continue ad infinitum! More practically, the process ends once the query has been fully expressed as a SQL statement.

Subqueries can appear when using the comparison predicate, the IN predicate and when quantifiers are used (not yet explained).

Comparison Predicate

14. Get the supplier numbers of suppliers who are located in the same city as Tapi.

```

SELECT      S.S#, SNAME
FROM        S
WHERE       S.CITY =
(SELECT      S.CITY
FROM        S
WHERE       SNAME = 'TAPI' )

```

The inner select (subquery) retrieves the city of the supplier named Tapi. The outer select (the main one) then compares the city of each supplier in the supplier relation and picks up those where the comparison succeeds.

```

S#          SNAME
11         NARMADA
13         TAPI

```

Notice that the subquery appears after the comparison operator. The format of this form of expression

scalar-expr operator subquery

In this form the subquery selects a set of values. The outer query checks whether the value of a specified attribute is in this set.

15. Get the names of suppliers who supply part 2

```
SELECT      S.SNAME
FROM        S
WHERE       S.S# IN
(SELECT      SP.S#
FROM        SP
WHERE       SP.P# = 2)
```

SNAME

CAUVERY

NARMADA

The above query can be equivalently expressed as

```
SELECT      S.SNAME
FROM        S
WHERE       2 IN
(SELECT      P#
FROM        SP
WHERE       S.S# = P#)
```

S# is unqualified and therefore, refers to SP. That is because every unqualified attribute name is implicitly qualified with the relation name from the nearest applicable FROM clause.

Quantified predicates

The two quantifiers that can be used are the ALL and ANY. Any stands for the existential quantifier and ALL for the universal quantifier.

Lets first look at ANY. It can be specified in a comparison predicate just following the comparison operator. That is,

scalar-expr O ANY subquery

The subquery is first evaluated to give a set of values. The above expression is true if the scalar-expr is O comparable with any of the values that form the result of the subquery.

16. Get the part numbers for parts whose cost is less than the current maximum cost.

```
SELECT      P#, COST
```

```
FROM        P
```

```
WHERE       COST < ANY
```

```
(SELECT      COST
```

```
FROM        P)
```

P#	COST
1	10
2	15

P#	COST
1	10
2	15

The inner select gets the cost of all the parts. In the outer select, a P# is selected if its cost is less than some element of the set selected in the earlier step.

17. Get the supplier names of suppliers who do not supply part 2.

```
SELECT      SNAME
FROM        S
WHERE       2 < > ALL
(SELECT      P#
FROM        SP
WHERE       S# = S.S#)
```

SNAME

YAMUNA

TAPI

For each supplier, all the parts supplied by him are collected in the inner select. If none of them is equal to P2 then the condition evaluates to true and supplier name forms part of the result.

Existence Test

This kind of an expression is used when it is necessary to find out if the set of values retrieved by using a subquery contains an element or not.

18. Get the supplier names of suppliers who supply at least one part.

```
SELECT      SNAME
FROM        S
WHERE       EXISTS
(SELECT      *
FROM        SP
WHERE       SP.S# = S.S#)
```

For a given supplier, if the subquery selects at least one tuple then the condition (which follows WHERE) evaluates to true. Then, the name of the supplier is selected. IN our data base every supplier is supplying at least one part. So the names of all of them would be part of the result.

3.4.3 Functions

Some standard functions are defined in SQL and can be used when framing queries. There are five built-in functions. These are

COUNT	-	number of values of a column
SUM	-	Sum of values in a column
AVG	-	Average of values in a column
MAX	-	Maximum of all the values in a column
MIN	-	Minimum of all the values in a column

If the function is followed by the word DISTINCT then unique values are used. On the other hand, if ALL follows the function then all the values are used for evaluating the function. ALL is the default.

COUNT(*) has a special meaning in that it counts the number of rows of a relation. COUNT

in any other form must make use of DISTINCT. In other words, except when rows are counted, COUNT always returns the number of distinct values in a column.

19. Get the total number of suppliers

```
SELECT COUNT(*)
FROM S
```

COUNT(*) counts the number of tuples of S and hence, the number of suppliers.

20. Get the total quantity of Part 2 that is supplied.

```
SELECT SUM(SP.QTY)
FROM SP
WHERE SP.P# = 2
```

The answer is one of

- a) 25 b) 40 c) 80 d) 100

21. Get the part numbers whose cost is greater than the average cost.

```
SELECT P#
FROM P
WHERE COST
  > (SELECT AVG(COST)
      FROM P)
```

22. Get the names of suppliers who supply from a city where there is atleast one more supplier.

```
SELECT SNAME
FROM S
WHERE 2
  <= (SELECT COUNT(CITY)
      FROM S
      WHERE CITY = FIRST.CITY)
```

Some practice exercise before we move on to the next section. Try and write the expression yourself before looking at the solution. Yours might be different from the solution given in the notes. For all you know, yours might be a better solution. So, go-ahead and try. Use any feature that has been covered till now.

1. Get the names of suppliers who supply at least one red part

```
SELECT SNAME
FROM S
WHERE S# IN
  (SELECT S#
   FROM SP
   WHERE P# IN
     (SELECT P#
      FROM P
      WHERE COLOUR = 'RED'))
```

2. Get the supplier numbers who supply at least one part supplied by supplier 10.

```

SELECT      DISTINCT S#
FROM        SP
WHERE       P# IN
(SELECT      P#
FROM        SP
WHERE       S# = 10 )

```

3.4.4 GROUP BY Feature

This feature allows one to partition the result into a number of groups such that all rows of the group have the same value in some specified column.

23. Get the part number and the total quantity.

```

SELECT      P#, SUM(QTY)
FROM        SP
GROUP BY    P#
P#          SUM(QTY)
1           125
2           80
3           110

```

GROUP BY groups together all the rows which have the same value for P#. The function SUM is then applied to each group. That is, the result consists of a part number along with the total quantity in which it is supplied.

Whenever GROUP BY is used then the phrase WHERE is to be replaced by HAVING. The meaning of HAVING is the same as WHERE except that the condition is now applicable to each group.

24. Get the part numbers for parts supplied by more than one supplier.

```

SELECT      P#
FROM        SP
GROUP BY    P#
HAVING     COUNT(*) > 1

```

Each group contains one or more tuples which have the same part number. COUNT(*) is applied to each such group.

The result before COUNT(*) is applied is

```

P#
1
2
3

```

In this case all the part numbers will be selected.

3.4.5 Updating the Database

The contents of the database can be modified by inserting a new tuple, deleting an existing tuple or changing the values of attributes of one or more tuples.

INSERT

Structured Query Language

The insertion facility allows new tuples to be inserted into given relations. Attributes which are not specified by the insertion statement are given null values. Consider

1. Add a part with number 14, weight 10, coloured red, with the cost and selling price as 20 and 60 respectively.

INSERT INTO P :

< 14, 10, 'red', 20, 60 >

The tuple is inserted into P.

If all the fields are not known then a tuple can still be added. The attributes whose values are not specified will have a null value.

INSERT INTO P :

<15, 'GREEN'>

The values for fields the weight, cost and the selling price which are not specified are assumed to be null.

2. Let us assume that there is a relation called RED-PART with one column P#.

INSERT INTO RED-PART :

```
SELECT    P#
FROM      P
WHERE     COLOUR = 'red'
```

The various attributes of P having red colour are identified and inserted into the relation RED-PART.

DELETE

The deletion facility removes specified tuples from the database. Consider

1. Delete supplier 13

```
DELETE    S
WHERE    S# = 13
```

Since S# is the primary key only one tuple will be deleted from S.

2. Delete all suppliers who supply from Bangalore

```
DELETE    S
WHERE    SCITY = 'BANGALORE'
```

Here, more than one supplier can get deleted.

3. Delete all the suppliers

```
DELETE    S
```

The definition of S exists but the relation is empty.

4. Delete all the supplies involving red coloured parts.

```
DELETE    SP
WHERE    'red' =
(SELECT    COLOUR
FROM      P
WHERE    P.P# = SP.P#)
```

UPDATE

When columns are to be modified SET clause is used. This clause specifies the update to be made to selected tuples.

1. Change the city of supplier 13 to Bangalore and increase the turnover by 20 lakhs.

```
UPDATE S
SET CITY = 'BANGALORE'
TURNOVER = TURNOVER + 20
WHERE S# = 13

2. Increase quantity by 10 for all supplies of red coloured parts.

UPDATE SP
SET QTY = QTY + 10
WHERE P# IN
(SELECT P#
FROM P
WHERE COLOUR = 'RED')
```

3.4.6 Data Definition Facilities

Data definition facilities permit users to create and drop relations, define alternative views of relations.

CREATE statement allows to define a relation. The name of the relation to be created and its various fields together with their data types must be specified. If a certain attribute is barred from containing null values then a NONNULL specification must be made for it.

It must be noted that the word TABLE is used in this syntax instead of RELATION.

Example

```
CREATE TABLE DEPT
(DNO(CHAR(2),NONULL),
DNAME (CHAR (12) VAR),
LOC(CHAR(20) VAR))
```

VIEW

A very important aspect of data definition is the ability to define alternative views of data. The process of specifying an alternative view is very similar to that of framing a query. The derived relation is stored and can be used thereafter as an object of the various commands. It is also possible to define other views on top of the newly created relation.

Example

```
DEFINE VIEW D50 AS
SELECT EMPNO, NAME, JOB
FROM EMP
WHERE DNO = 50
```

D50 contains the employee number, name and job of those employees who are in department 50.

3.5 VIEWS

A view is a virtual table, that is one that does not actually exist. It is made up of a query on other tables in the database. It could include only certain columns or rows from a table or from many tables. A view which restricts the user to certain rows is called a horizontal view and a vertical view restricts the user to certain columns. You are not restricted to purely horizontal or vertical slices of data.

A view can be as complicated as you like. You can have grouped views where the query contains a GROUP BY clause. This makes the view a summary of the data in a table or tables.

If the list of column names is omitted the columns in the view take the same name as in the underlying tables. You must specify column names if the query includes calculated columns or two columns with the same name. There are several advantages to views, including :

- **Security** : Users can be given access to only the rows/columns of tables that concern them. The table may contain data for the whole firm but they only see data for their department.
- **Date integrity** : The WITH CHECK OPTION clause is used to enforce the query conditions on any updates to the view. If the view shows data for a particular office the user can only enter data for that office.
- **Simplicity** : Even if a view is a multi-table query, querying the view still looks like a single-table query.
- **Protection from change** : If the structure of the database changes, the user's view of the data can remain the same.

There are two disadvantages to views :

- **Performance** : A view may look like a single table but underneath the DBMS is usually still running multi-table queries. If the view is complex then even simple queries can take a long time.
- **Update restrictions** : Updating the data through a view may or may not be possible. If the view is complex the DBMS may decide it can't perform updates and make the view read-only.

The ISO standard specifies five conditions that a view must meet in order to allow updates :

- The view must not have a DISTINCT clause
- The view must only name one table in the FROM clause
- All columns must be real columns — no expressions, calculated columns or column functions
- The WHERE clause must not contain a sub-query
- There must be no GROUP BY or HAVING clause

You will find that most dialects of SQL are not quite so restrictive. The underlying principle is that updates are allowed if the rows and columns of the view are traceable back to actual rows and columns in tables.

The format of view statement is as follows :

create view <view name> as query expression

A view is a relation (virtual rather than base) and can be used in query expressions, that is, queries can be written using the view as a relation. Views generally are not stored, since the data in the base relations may change. The base relations on which a view is based are sometimes called the **existing relations**. The definition of a view in a **create view** statement is stored in the system catalog. Having been defined, it can be used as if the view really represented a real relation. However, such a virtual relation defined by a view is recomputed whenever a query refers to it.

Example

- (a) For reasons of confidentiality, not all users are permitted to see the Pay_Rate of an employee. For such user the DBA can create a view, for example, EMP_VIEW defined as:

```
create view EMP_VIEW as
(select Empl_No, Name, Skill
from EMPLOYEE)
```

- (b) A view can be created for a subset of the tuples of a relation, as in this example. For assigning employees to particular jobs, the manager requires a list of the employees who have not been assigned to any jobs:

```
create view FREE as
(select Empl_No
from EMPLOYEE)
minus
(select Empl_No
from DUTY_ALLOCATION)
```

- (c) The view in part (b) above can also be created using the following statements:

```
create view FREE as
(select Empl_No
from EMPLOYEE)
where Empl_No any
(select Empl_No
from DUTY_ALLOCATION)
```

In the above examples, the names of the attributes in the views are implicitly taken from the base relation. The data types of the attribute of the view are inherited from the corresponding attributes in the base relation. We can, however, give new names to the attributes of the view. This is illustrated in the syntax of the create view statement given below:

```
create view VIEW_NAME
(Name1, Name 2, ....)
as (select .....)
```

Here the attributes in the view are given as Name1, Name 2, and these names are associated with the existing relation by order correspondence. The definition of a view is accomplished by means of a subquery involving a select statement as given in the syntax above. Since a view can be used in a select statement, a view can be defined on another existing view.

3.6 SUMMARY

Most commercial relational DBMSs support some form of the SQL data manipulation language, and this creates different dialects of SQL. SQL has been standardised; that is, a minimum compatible subset is specified as a standard. In addition, embedded versions of SQL are supported by many commercial DBMSs. This allows application program written in a high-level language such as BASIC, C, COBOL, FORTRAN, Pascal, or PL/I to use the database accessing SQL by means of appropriate preprocessors.

3.7 FURTHER READING

Bipin C. Desai, *An Introduction to Database Management*, Golgotia Publication, New Delhi.