

UNIT 1 INTRODUCTION TO OBJECT ORIENTED DATABASE MANAGEMENT SYSTEM

Structure

- 1.0 Introduction
 - 1.1 Objectives
 - 1.2 What are Next Generation Data Base System?
 - 1.3 New Database Application
 - 1.4 What is Object Oriented Database Management System?
 - 1.5 Promises of Object Oriented System
 - 1.6 Promises and Advantages of Object Oriented Database Management System
 - 1.7 Deficiencies of Relational Database Management System
 - 1.8 Difference Between Relational Database Management System and Object Oriented Database Management System
 - 1.9 Alternative Objective Oriented Database Strategies
 - 1.10 Summary
 - 1.11 Model Answers
 - 1.12 Further Readings
-

1.0 INTRODUCTION

Since 1960s, Data Base Management Systems (DBMS) have been widely used in data processing environment. The support of characteristics such as data sharing, independence, consistency, integrity is the main reason for its success which traditional file management system does not inherently offer.

A database system is usually organised according to a data model. In the previous block, we discussed three most popular models: hierarchical, network and relational. The difference among all these three models is in the way of organising records, although they are record based. They were mainly designed to process large amount of relatively simple and fixed format data. DBMS based on these models along with sophisticated indexing and query optimization techniques have served business oriented database application especially well.

RDBMSs were originally designed for mainframe computer and business data processing applications. Moreover, relational systems were optimized for environments with large number of users who issue short queries. But today's application has moved from centralised mainframe computer to networked workstation on every desk. These applications include computer aided design (CAD), multimedia system, software engineering (design of complex project), knowledge database (to be discussed in unit 3 of this block). These operations require complex operations and data structure representation. For example, a multimedia database may contain variable length text, graphics, images, audio and video data. Finally a knowledge base system requires data rich in semantics.

Existing commercial DBMS, both small and large scale have proven inadequate for these applications. The traditional database notion of storing data in two-dimensional tables or in flat files breaks down quickly in the face of complex data structures and data types used in today's applications.

Research to model and process complex data has gone in two directions:

- (a) extending the functionality of RDBMS
- (b) developing and implementing OODBMS that is based on object oriented programming paradigm.

OODBMSs are designed for use in today's application areas such as multimedia, CAD, office automation, etc. In this unit, we will touch up some of the basic issues related to OODBMS.



1.1 OBJECTIVES

After going through this unit, you will be able to:

- define what is object oriented DBMS
- differentiate between RDBMS and OODBMS
- list next generation database systems
- list advantages of object oriented DBMS

1.2 WHAT ARE NEXT GENERATION DATABASE SYSTEM?

Computer sciences have gone through several generation of database management starting with indexed files and later, network and hierarchical data base management systems (DBMS). More recently, relational DBMS revolutionised the industry by providing powerful data management capabilities based on few simple concepts. Now, we are on the verge of another generation of database system called **Object Oriented DBMS** based on object oriented programming paradigm. This new kind of DBMS, unlike previous DBMS models, manages more complex kind of data for example multimedia objects. The other kind of next generation DBMS is knowledge database management system (KDBMS) which is used to support the management of the shared knowledge. It supports a large number of complex rules for automatic data inferencing (retrieval) and maintenance of data integrity.

The goal of these new DBMS is to support a much wider range of data intensive applications in engineering, graphic representation-scientific and medical. These new DBMS can also support new generations of traditional business applications.

1.3 NEW DATABASE APPLICATIONS

Some applications that require the manipulation of large amounts of data can benefit from using a DBMS. However, the nature of the data in these applications does not fit well into the relational framework.

- (1) Design databases: Engineering design databases are useful in computer-aided design/manufacturing/software engineering (CAD/CAM/CASE) systems. In such systems, complex objects can be recursively partitioned into smaller objects. Furthermore, an object can have different representations at different levels of abstraction (equivalent objects). Moreover, a record of n object's evolution (object versions) should be maintained. Traditional database technology does not support the notions of complex objects, equivalent objects, or object versions.
- (2) Multimedia databases: In a modern office information or other multi-media system, data include not only text and numbers but also images, graphics and digital audio and video. Such multimedia data are typically stored as sequences of bytes with variable lengths, and segments of data are linked together for easy reference. The variable length data structure cannot fit well into the relational framework, which mainly deals with fixed-format records. Furthermore, applications may require access to multimedia data on the basis of the structure of a graphical item or by following logical links. Conventional query languages were not designed for such applications.
- (3) Knowledge bases: Artificial intelligence and expert systems represent; information as facts and rules that can be collectively viewed as a knowledge base. In typical Artificial Intelligence applications, knowledge representation requires data structures with rich semantics that go beyond the simple structure of the relational model. Artificial decomposition and mapping would be necessary if a relational DBMS were used. Furthermore, operations in a knowledge base are more complex than those in a traditional database. When a rule is added, the system must check for contradiction and redundancy. Such operations cannot be represented directly by relational operations, and the complexity of checking increases rapidly as the size of the knowledge base grows.



In general, these applications require the representation of complex data elements as well as complex relationships among them. Users in these environments have found relational technology inadequate in terms of flexibility, modeling power, and efficiency.

1.4 WHAT IS OBJECT ORIENTED DATABASE MANAGEMENT SYSTEM?

Object-oriented technologies in use today include object-oriented programming languages (e.g., C++ and Smalltalk), object-oriented database systems, object-oriented user interfaces (e.g., Macintosh and Microsoft Windows systems) and so on. An object-oriented technology is a technology that makes available to the users facilities that are based on object-oriented concepts. To define object-oriented concepts, we must first understand what an **object** is.

Object

The term object means a combination of data and program that represents some real-world entity. For example, consider an employee named Amit; Amit is 25 years old, and his salary is \$25,000. Then Amit may be represented in a computer program as an object. The data part of this object would be (name: Amit, age: 25, salary: \$25,000). The program part of the object may be a collection of programs (hire, retrieve the data, change age, change salary, fire). The data part consists of data of any type- For the Amit object, string is used for the name, integer for age; and monetary for salary; but in general, even any user-defined type, such as Employee, may be used. In the Amit object, the name, age, and salary are called attributes of the object.

Encapsulation

Often, an object is said to **encapsulate** data and program. **This means that the users cannot see the inside the object but can use the object by calling the program part of the object.** This is not much different from procedure calls in conventional programming; the users call a procedure by supplying values for input parameters and receive results in output parameters.

Inheritance and Class

The term **object-oriented** roughly means a combination of object encapsulation and inheritance. The term inheritance is sometimes called reuse. Inheritance means roughly that a new object may be created by extending an existing object. Now let us understand the term inheritance more precisely. An object has a data part and a program part. **All objects that have the same attributes for the data part and same program part are collectively called a class (or type).** The classes are arranged such that some class may inherit the attributes and program part from some other classes.

Amit, Ankit and Anup are each an Employee object. The data part of each of these objects consists of the attributes Name, Age and salary. Each of these Employee objects has the same program part (hire, retrieve the data, change age, change salary, fire). **Each program in the program part is called a method.** The term class refers to the collection of all objects that have the same attributes and methods. In our example, the Amit, Ankit and Anup objects belong to the class Employee since they all have the same attributes and methods. This class may be used as the type of an attribute of any object. At this time, there is only one class in the system namely, the class Employee; and three objects that belong to the class namely Amit, Ankit and Anup objects.

Inheritance Hierarchy or Class Hierarchy

Now suppose that a user wishes to create two sales employees, Jai and Prakash. But sales employees have an additional attribute namely, **commission**. The sales employees cannot belong to the class Employee. However, the user can create a new class, Sales Employee, such that all attributes and methods associated with the class Employee may be reused and the attribute commission may be added to Sales Employee. The user does this by declaring the class Sales Employee to be a subclass of the class Employee. The user can now proceed to create the two sales employees as objects belonging to the class Sales Employee. The users can create new classes as subclasses of existing classes. **In general, a class may inherit from one or more existing classes** and the inheritance structure of classes becomes a directed acyclic graph (DAG); but for simplicity, the inheritance structure is called an **inheritance hierarchy or class hierarchy**.



The power of object-oriented concepts is delivered when encapsulation and inheritance work together.

- Since inheritance makes it possible for different classes to share the same set of attributes and methods, the same program can be run against objects that belong to different classes. This is the basis of the object-oriented user interface that desktop publishing systems and windows management systems provide today. The same set of programs (e.g., open, close, drop, create, move, etc.) apply to different types of data (image, text file, audio, directory, etc.).
- If the users define many classes, and each class has many attributes and methods, the benefit of sharing not only the attributes but also the programs can be dramatic. The attributes and programs need not be defined and written from scratch. New classes can be created by adding attributes and methods of existing classes, thereby reducing the opportunity to introduce new errors to existing classes.

1.5 PROMISES OF OBJECT ORIENTED SYSTEMS

Object-oriented systems make these promises:

- **Reduced maintenance**

The primary goal of object-oriented development is the assurance that the system will enjoy a longer life while having far smaller maintenance costs. Because most of the processes within the system are encapsulated, the behaviours may be reused and incorporated into new behaviours.

- **Real-world modeling**

Object-oriented systems tend to model the real world in a more complete fashion than do traditional methods. Objects are organised into classes of objects, and objects are associated with behaviours. The model is based on objects rather than on data and processing.

- **Improved reliability**

Object-oriented systems promise to be far more reliable than traditional systems, primarily because new behaviours can be built from existing objects.

- **High code reusability**

When a new object is created, it will automatically inherit the data attributes and characteristics of the class from which it was spawned. The new object will also inherit the data and behaviours from all superclasses in which it participates.

1.6 PROMISES AND ADVANTAGES OF OBJECT ORIENTED DATABASE MANAGEMENT SYSTEM

An object-oriented programming language (OOPL) provides facilities to create classes for organising objects, to create objects, to structure an inheritance hierarchy to organise classes so that subclasses may inherit attributes and methods from superclasses, and to call methods to access specific objects. Similarly, an object-oriented database system (OODB) should provide facilities to create classes for organising objects, to create objects, to structure an inheritance hierarchy to organise classes so that subclasses may inherit attributes and methods from superclasses, and to call methods to access specific objects. Beyond these, an OODB, because it is a database system, must provide standard database facilities found in today's relational database systems (RDBs), including nonprocedural query facility for retrieving objects, automatic query optimisation and processing, **dynamic schema changes** (changing the class definitions and inheritance structure), **automatic management of access methods** (e.g., B+-tree index, extensible hashing, sorting, etc.) to improve **query processing performance, automatic transaction management, concurrency control, recovery from system crashes, and security and authorisation**, Programming languages, including OOPLs, are designed with one user and a relatively small database in mind. Database systems are designed with many users and very large databases in mind; hence performance, security and authorisation, concurrency control, and dynamic schema changes become important issues. Further, transaction systems are used to maintain critical data accurately; hence, transaction management, concurrency control, and recovery are important facilities.



In so far as a database system is a system software, whose Functions are called from application programs written in some host programming languages, we may distinguish two different approaches to designing an OODB. One is to store and manage objects created by programs written in an OOPL. Some of the current OODBs are designed to store and manage objects generated in C++ or Smalltalk programs. Of course, an RDB can be used to store and manage such objects. However, RDBs do not understand objects-in particular, methods and inheritance. Therefore, what may be called an object manager or an object-oriented layer software needs to be written to manage methods and inheritance and to translate objects to tuples (rows) of a relation (table). But the object manager and RDB combined are in effect an OODB (with poor performance, of course).

Another approach is to make object-oriented facilities available to users of non-OOPLs. The users may create classes, objects, inheritance hierarchy, and so on, and the database system will store and manage those objects and classes. This approach in effect turns non-OOPLs (e.g., C, FORTRAN, COBOL, etc.) into object-oriented languages. In fact, C++ has turned C into an OOPL, and CLOS has added object-oriented programming facilities to Common LISP. An OODB designed using this approach can of course be used to store and manage objects created by programs written in an OOPL. Although a translation layer would need to be written to map the OOPL objects to objects of the database system, the layer should be much less complicated than the object manager layer that an RDB would require.

In view of the fact that C++, despite its growing popularity, is not the only programming language that database application programmer's are using or will ever use, and there is a significant gulf between a programming language and a database system that will deliver the power of object-oriented concepts to database application programmers. Regardless of the approach, OODBs, if done right, can bring about a quantum jump in the productivity of database application programmers and even in the performance of the application programs.

One source of the technological quantum jump is the reuse of a database design and program that object-oriented concepts make possible for the first time in the evolving history of database technologies. Object-oriented concepts are fundamentally designed to reduce the difficulty of developing and evolving complex software systems or designs. Encapsulation and inheritance allow attributes (i.e., database design) and programs to be reused as the basis for building complex databases and programs. This is precisely the goal that has driven the data management technology from file systems to relational database systems during the past three decades. An OODB has the potential to satisfy the objective of reducing the difficulty of designing and evolving very large and complex databases.

Another source of the technological jump is the powerful **data type facilities** implicit in the object-oriented concepts of encapsulation and inheritance.

Advantages of Object-Oriented Databases

Systems developed with object-oriented languages have many benefits, as previously discussed. Yet, as also described, these systems have particular attributes that can be complemented with object-oriented databases. These attributes include lack of persistence, inability to share objects among multiple users, limited version control, and lack of access to other data, for example, data in other databases.

In systems designed with object-oriented languages, objects are created during the running of a program and are destroyed when the program ends. Providing a database that can store the objects between runs of a program offers both increased flexibility and increased security. The ability to store the objects also allows the objects to be shared in a distributed environment. An object-oriented database can allow only the actively used objects to be loaded into memory and thus minimizes or preempts the need for virtual memory paging. This is especially useful in large-scale systems.

Persistent objects also allow objects to be stored for each version. This version control is useful not only for testing applications, but also for many object-oriented design applications where version control is a functional requirement of the application itself. Access to other data sources can also be facilitated with object-oriented databases, especially those built as hybrid relational systems, which can access relational tables as well as other object types.

Object-oriented databases also offer many of the benefits that were formerly found only in expert systems. With an object-oriented database, the relationships between objects and the constraints on objects are maintained by the database management system, that is, the objects themselves. The rules associated with the expert system are essentially replaced by the object schema and the methods. As many expert systems currently do not have adequate



database support, object-oriented databases afford the possibility of offering expert system functionality with much better performance.

Object-oriented databases offer benefits over current hierarchical and relational database models. They enable support of complex applications not supported well by the other models. They enhance programmability and performance, improve navigational access, and simplify concurrency control. They lower the risks associated with referential integrity, and they provide a better user metaphor than the relational model.

Object-oriented databases by definition allow the inclusion of more of the code (i.e. the object's methods) in the database itself. This incremental knowledge about the application has a number of potential benefits of the database system itself, including the ability to optimize query processing and to control the concurrent execution of transactions.

Performance, always a significant issue in system implementation, may be significantly improved by using an object-oriented model instead of a relational model. The greatest improvement can be expected in applications with high data complexity and large numbers of inter-relationships. Clustering, or locating the related objects in close proximity, can be accomplished through the class hierarchy or by other interrelations. Caching, or the retention of certain objects in memory or storage, can be optimised by anticipating that the user or application may retrieve a particular instance of the class. When there is high data complexity, clustering and caching techniques in object databases gain tremendous performance benefits that relational databases, because of their fundamental architecture, will never be able to approach.

Object-oriented databases can store not only complex application components but also larger structures. Although relational systems can support a large number of tuples (i.e. rows in a table), individual types are limited in size. Object-oriented databases with large objects do not suffer a performance degradation because the objects do not need to be broken apart and reassembled by applications,' regardless of the complexity of the properties of the application objects.

Since objects contain direct references to other objects, complex data set can be efficiently assembled using these direct references. The ability to search by direct references significantly improves navigational access. In contrast, complex data sets in relational databases must be assembled by the application program using the slow process of joining tables.

For the programmer, one of the challenges in building a database is the data manipulation language (DML) of the database. DMLs for relational databases usually differ from the programming language used to construct the rest of the application. This contrast is due to differences in the programming paradigms and mismatches of type systems. The programmer must learn two languages, two tool sets, and two paradigms because neither alone has the functionality to build an entire application. Certain types of programming tools such as application generators and fourth-generation languages (4GLs) have emerged to produce code for the entire application, thereby bridging the mismatch between the programming language and the DML, but most of these tools compromise the application programming process.

With object-oriented databases much of this problem is eliminated. The DML can be extended so that more of the application can be "written in the DML. Or an object-oriented application language, of example C++ can be extended Lo be the DML. More or the application can be built into the database itself. Movement across the programming interface between the databases the application then occurs in a single paradigm with a common set of tools, Class libraries can also assist the programmer in speeding the creation of databases. Class libraries encourage reuse of existing code and help to minimise the cost of later modifications. Programming is easier because the data structures model the problem more closely. Having the data and procedures encapsulated in a single object makes it less likely that a change to one object will affect the integrity of other objects in the database.

Concurrency control is also simplified with an object-oriented database. In a relational database, the application needs to lock each record in each table explicitly because related data re-represented across a number of tables. Integrity, a key requirement for databases, can be better supported with an object-oriented database, because the application can lock all the relevant data in one operation. Referential integrity is better supported in an object-oriented database because the pointers are maintained and updated by the database itself. Finally, object-oriented databases offer a better user metaphor than relational databases. The tuple or table, although enabling a well-defined implementation strategy, is not an intuitive modeling



framework, especially outside the domain of numbers. Objects offer a more natural and encompassing modeling metaphor.

1.7 DEFICIENCIES OF RELATIONAL DATA BASE MANAGEMENT SYSTEM

The data type facilities in fact are the keys to eliminating three of the important deficiencies of RDBs. These are summarized below, we will discuss these points in greater detail later.

- RDBs force the users to represent hierarchical data (or complex nested data or compound data) such as bill of materials in terms of tuples in multiple relations. This is awkward to start with. Further, to retrieve data thus spread out in multiple relations. RDBs must resort to joins, a generally expensive operation. The data type of an attribute of an object in OOPLs may be a primitive type or an arbitrary user-defined type (class). The fact that an object may have an attribute whose value may be another object naturally leads to nested object representation, which in turn allow hierarchical data to be naturally (i.e., hierarchically) represented.
- RDBs offer a set of primitive built-in data types for use as domains of columns of relation, but they do not offer any means of adding user-defined data types. The built-in data types are basically all numbers and short symbols. RDBs are not designed to allow new data types to be added and therefore often require major surgery to the system architecture and code to add any new data type. Adding a new data type to a database system means allowing its use as the data type of an attribute—that is, storage of data of that type, querying, and updating of such data. Object encapsulation in OOPLs does not impose any restriction on the types of data that the data may be primitive types or user-defined types. Further, new data types may be created as new classes, possibly even as subclasses of existing classes, inheriting their attributes and methods.
- Object encapsulation is the basis for the storage and management of programs as well as data in the database. RDBs now support stored procedures—that is, they allow programs to be written in some procedural language and stored in the database for later loading and execution. However, the stored procedures in RDBs are not encapsulated with data—that is, they are not associated with any relation or any tuple of a relation. Further, since RDBs do not have the inheritance mechanism, the stored procedures cannot automatically be reused.

1.8 DIFFERENCE BETWEEN RELATIONAL DATABASE MANAGEMENT SYSTEM AND OBJECT ORIENTED DATABASE MANAGEMENT SYSTEM

RDBMSs were never designed to allow for the nested structure. These types of applications are extensively found in CAD/CAE, aerospace, etc. OODBM can easily support these applications. Moreover, it is much easier and natural to navigate through these complex structures in form of objects that model the real world in OODBMS rather than table, tuples and records in RDBMS. It is hard to confuse a relational database with an object-oriented database. The normalised relational model is based on a fairly elegant mathematical theory. Relational databases derive a virtual structure at run time based on values from sets of data stored in tables. Databases construct views of the data by selecting data from multiple tables and loading it into a single table (OODBs traverse the data from object to object).

Relational databases have a limited number of simple, built-in data types, such as **integer** and **string**, and a limited number of built-in operations that can handle these data types. You can create complex data types in a relational database, but you must do it on a linear basis, such as combining fields into records. And the operations on these new complex types are restricted, again, to those defined for the basic types (as opposed to arbitrary data types or subclassing with inheritance as found in OODBs).



than the reinvention, of commonly used data elements. Objects in an OODB survive multiple sessions; they are persistent. If you delete an object stored in a relational database, other objects may be left with references to the deleted one and may now be incorrect. The integrity of the data thus becomes suspect and creates inconsistent versions.

In the relational database, complex objects must be broken up and stored in separate tables. This can only be done in a sequential procedure with the next retrieval relying on the outcome of the previous. The relational database does not understand a global request and thus cannot optimise multiple requests; OQDBs can issue a single message (request) that contains multiple transactions.

The relational model, however, suffers at least one major disadvantage. It is difficult to express the semantics of complex objects with only a table model for data storage. Although relational databases are adequate for accounting or other typical transaction-processing applications where the data types are simple and few in number, the relational model offers limited help when data types become numerous and complex.

Object-oriented databases are favored for applications where the relationships among elements in the database carry the key information. Relational databases are favored when the values of the database elements carry the key information. That is, object-oriented models capture the structure of the data; relational models organise the data itself. If a record can be understood in isolation, then the relational database is probably suitable. If a record makes sense only in the context of other records, then an object-oriented database is more appropriate.

Engineering and technical applications were the first applications to require databases that handle complex data types and capture the structure of the data. Applications such as mechanical and electrical computer-aided design (MCAD and ECAD) have always used nontraditional forms of data, representing such phenomena as three-dimensional images and VLSI circuit designs. Currently these application programs store their data in application-specific file structures. The data-intensiveness of these applications is not only in the large amount of data that need to be programmed into the database, but in the complexity of the data itself. In these design-based applications, relationships among elements in the database carry key information for the user. Functional requirements for complex cross references, structural dependences, and version management all require a richer representation than what is provided by hierarchical or relational databases.

Check Your Progress

1. What are the drawbacks of current commercial databases?

.....
.....
.....

2. What is the meaning of multi-media data?

.....
.....
.....

3. List few requirements for multi-media data management.

.....
.....
.....



1.9 ALTERNATIVE OBJECT-ORIENTED DATABASE STRATEGIES

There are at least six approaches for incorporating object orientation capabilities in databases:

1. **Novel database data model/data language approach:** The most aggressive approach is to develop entirely new database language and database management system with object orientation capabilities. Most of the research projects in object-oriented databases have pursued this approach. In the industry introduces novel DML (Data Manipulation Language) and DDL (Data Definition Language) constructs for a data model based on semantic and functional data models.
2. **Extending an existing database language with object orientation capabilities:** A number of programming languages have been extended with object-oriented constructs. C++ flavors (an extension of LISP), and Object Pascal are examples of this approach in programming languages. It is conceivable to follow a similar strategy with database languages. Since SQL is a standard and the most popular database language, the most reasonable solution is to extend this language with object-oriented constructs, reflecting the object orientation capabilities of underlying database management system. This approach being pursued by most vendors of relational systems, as they evolve the next generation products. There have been many such attempts incorporating inheritance, function composition for nested entities, and even some support of encapsulation in an SQL framework.
3. **Extending on existing object-oriented programming language with database capabilities:** Another approach is to introduce database capabilities to an existing object-oriented language. The object orientation features abstract data typing, inheritance, object identity—will already be supported by the object-oriented language. The extensions will incorporate database features (querying, transaction support, persistence, and so on).
4. **Embedding object-oriented database language constructs in a host (conventional) language:** Database languages can be embedded in host programming languages. For example, SQL statements can be embedded in PL/1 C, FORTRAN and Ada. The types of SQL (that is relations and rows in relations) are quite different from the type systems of these host languages. Some object-oriented databases have taken a similar approach with a host language and an object-oriented database language.

1.10 SUMMARY

During the past decade, object oriented technology has found its way into database user interface, operating system, programming languages, expert system and the like. Object Oriented database product is already in the market for several years and several vendors of RDBMS are now declaring that they will extend their products with object oriented capabilities. In spite of all these claims there is no wide acceptability of OODBMS because of lack of industry standard. This technology is still evolving and takes some more time to get fully settled.

1.11 MODEL ANSWERS

Check Your Progress

Most of the current commercial database systems suffer from an inability to manage arbitrary types of data, arbitrary large data and data stored on devices other than magnetic disks. They understand a relatively limited set of data types such as integer, real data, monetary unit, short strings. Further they are not designed to manage data stored on such increasingly important storage devices such as CD-ROM and Videodisks.

Broadly, multimedia data means arbitrary data types and data from arbitrary data sources. Arbitrary data types include the numeric data and short string data supported in conventional database systems; large unstructured data, such as charts, graphs, tables, and arrays; and compound documents that are comprised of such data. Arbitrary data



sources include a native database; external (remote) databases; host file base; data input, storage, and presentation (output) devices; and even data generating and data-consuming programs (such as a text processing system).

- 3a) The ability to represent arbitrary data types (including compound documents) and specification of procedures (programs) that interact with arbitrary data sources.
- 3b) The ability to query, update, insert and delete multimedia data (including retrieval of multimedia data via associative search within multimedia data; minimally, text).
- 3c) The ability to specify and execute abstract operations on multimedia data; for example, to play, fast forward, pause, and rewind such one-dimensional data as audio and text; to display, expand and condense such two-dimensional data as a bit-mapped image.
- 3d) The ability to deal with heterogeneous data sources in a uniform manner; this includes access to data in these sources and migration of data from one data source to another.

1.12 FURTHER READINGS

1. Modern Database Systems—the Object Model, Interoperability and Beyond, By WON KIM, Addison Wesley, 1995.
2. Object-Oriented DBMS: Evolution & Performance Issues, A.R.Hurson & Simin H. Pakzad, IEEE Computer, Feb. 1993.